

Principles of Database Systems

Assignment #4 - Structured Query Language 3

He Tianyang, 3022001441

October 9, 2024

1 Execute the SQL in Slides

1.1 Preparation

1.1.1 Create Table

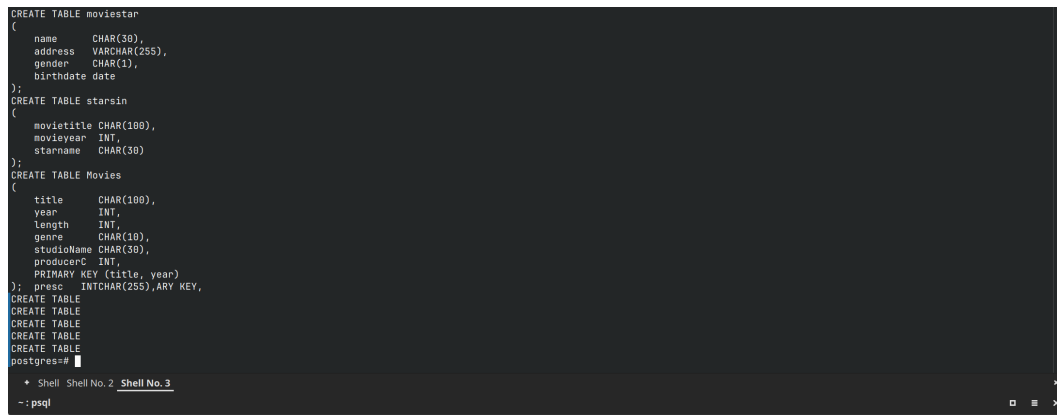
```
CREATE TABLE movieexec
(
    name      CHAR(30),
    address   VARCHAR(255),
    cert      INT,
    networth  INT
);
CREATE TABLE moviestar
(
    name      CHAR(30),
    address   VARCHAR(255),
    gender    CHAR(1),
    birthdate date
);
CREATE TABLE starsin
(
    movietitle CHAR(100),
    movieyear  INT,
    starname   CHAR(30)
);
CREATE TABLE Movies
(
    title     CHAR(100),
    year      INT,
```

```

length      INT,
genre       CHAR(10),
studioName  CHAR(30),
producerC   INT,
PRIMARY KEY (title, year)
);
CREATE TABLE studio
(
    name      CHAR(50) PRIMARY KEY,
    address   VARCHAR(255),
    presc     INT
);

```

the execution results are shown in Figure 1.



```

CREATE TABLE moviestar
(
    name      CHAR(30),
    address   VARCHAR(255),
    gender     CHAR(1),
    birthdate  date
);
CREATE TABLE starsin
(
    movietitle CHAR(100),
    movieyear  INT,
    starname   CHAR(30)
);
CREATE TABLE Movies
(
    title      CHAR(100),
    year       INT,
    length     INT,
    genre      CHAR(10),
    studioName CHAR(30),
    producerC  INT,
    PRIMARY KEY (title, year)
); presc INTCHAR(255),ARY KEY,
CREATE TABLE
CREATE TABLE
CREATE TABLE
CREATE TABLE
CREATE TABLE
postgres=#
+ Shell Shell No. 2 Shell No. 3
~: psql

```

Figure 1: Create Table

1.1.2 Insert Data

```

INSERT INTO movies VALUES ('Logan's run', 1976, NULL, 'sciFi', 'MGM',
123);
INSERT INTO movies VALUES ('Star Wars', 1977, 124, 'sciFi', 'Fox', 555);
INSERT INTO movies VALUES ('Empire Strikes Back', 1980, 111, 'fantasy',
'Fox', 555);
INSERT INTO movies VALUES ('Star Trek', 1979, 132, 'sciFi', 'Paramount',
345);
INSERT INTO movies VALUES ('Star Trek: Nemesis', 2002, 116, 'sciFi', '
Paramount', 345);
INSERT INTO movies VALUES ('Terms of Endearment', 1983, 132, 'romance',
'MGM', 123);
INSERT INTO movies VALUES ('The Usual Suspects', 1995, 106, 'crime', '
MGM', 456);
INSERT INTO movies VALUES ('Gone With the Wind', 1938, 238, 'drama', '
MGM', 123);

```

```
INSERT INTO movies VALUES ('Wayne''s World', 1992, 95, 'comedy', '
    Paramount', 123);
INSERT INTO movies VALUES ('King Kong', 2005, 187, 'drama', 'Universal',
    789);
INSERT INTO movies VALUES ('King Kong', 1976, 134, 'drama', 'Paramount',
    666);
INSERT INTO movies VALUES ('King Kong', 1933, 100, 'drama', 'Universal',
    345);
INSERT INTO movies VALUES ('Pretty Woman', 1990, 119, 'comedy', 'Disney'
    , 999);

INSERT INTO movieexec VALUES ('George Lucas', 'Oak Rd.', 555, 200000000)
;
INSERT INTO movieexec VALUES ('Ted Turner', 'Turner Av.', 333,
    125000000);
INSERT INTO movieexec VALUES ('Stephen Spielberg', '123 ET road', 222,
    100000000);
INSERT INTO movieexec VALUES ('Merv Griffin', 'Riot Rd.', 199,
    112000000);
INSERT INTO movieexec VALUES ('Calvin Coolidge', 'Fast Lane', 123,
    20000000);
INSERT INTO movieexec VALUES ('Garry Marshall', 'First Street', 999,
    50000000);
INSERT INTO movieexec VALUES ('J.J. Abrams', 'High Road', 345, 45000000)
;
INSERT INTO movieexec VALUES ('Bryan Singer', 'Downtown', 456, 70000000)
;
INSERT INTO movieexec VALUES ('George Roy Hill', 'Baldwin Av.', 789,
    20000000);
INSERT INTO movieexec VALUES ('Dino De Laurentiis', ' Beverly Hills',
    666, 120000000);
INSERT INTO movieexec VALUES ('AAA', ' Beverly Hills', 666, 120000000);

INSERT INTO moviestar VALUES ('Jane Fonda', 'Turner Av.', 'F', '
    1977-07-07');
INSERT INTO moviestar VALUES ('Alec Baldwin', 'Baldwin Av.', 'M', '
    1977-06-07');
INSERT INTO moviestar VALUES ('Kim Basinger', 'Baldwin Av.', 'F', '
    1979-05-07');
INSERT INTO moviestar VALUES ('Harrison Ford', 'Beverly Hills', 'M', '
    1977-07-07');
INSERT INTO moviestar VALUES ('Carrie Fisher', '123 Maple St.', 'F', '
    1999-09-09');
INSERT INTO moviestar VALUES ('Mark Hamill', '456 Oak Rd.', 'M', '
    1988-08-08');
INSERT INTO moviestar VALUES ('Debra Winger', 'A way', 'F', '1978-05-06'
    );
INSERT INTO moviestar VALUES ('Jack Nicholson', 'X path', 'M', '
    1949-05-05');
```

```
INSERT INTO moviestar VALUES ('Kevin Spacey', 'New York Av.', 'F', '1937-12-21');
INSERT INTO moviestar VALUES ('AAA', 'New York Av.', 'F', '1937-12-21');

INSERT INTO starsin VALUES ('Star Wars', 1977, 'Carrie Fisher');
INSERT INTO starsin VALUES ('Star Wars', 1977, 'Mark Hamill');
INSERT INTO starsin VALUES ('Star Wars', 1977, 'Harrison Ford');
INSERT INTO starsin VALUES ('Empire Strikes Back', 1980, 'Harrison Ford');
INSERT INTO starsin VALUES ('The Usual Suspects', 1995, 'Kevin Spacey');
INSERT INTO starsin VALUES ('Terms of Endearment', 1983, 'Debra Winger');
INSERT INTO starsin VALUES ('Terms of Endearment', 1983, 'Jack Nicholson');

INSERT INTO studio VALUES ('MGM', 'MGM Boulevard', 123);
INSERT INTO studio VALUES ('Fox', 'Hollywood', 555);
INSERT INTO studio VALUES ('Disney', 'Buena Vista', 999);
INSERT INTO studio VALUES ('Paramount', 'Hollywood', 345);
INSERT INTO studio VALUES ('Universal', 'Hollywood', 789);
```

1.2 Aggregation Operators

```
SELECT AVG(netWorth) FROM MovieExec;

SELECT SUM(netWorth) FROM MovieExec;

SELECT MIN(netWorth) FROM MovieExec;

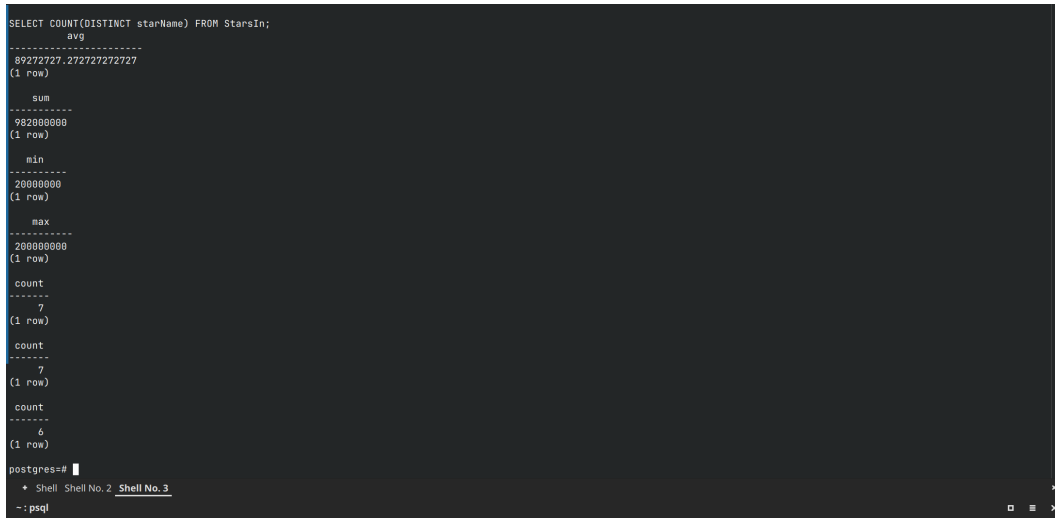
SELECT MAX(netWorth) FROM MovieExec;

SELECT COUNT(*) FROM StarsIn;

SELECT COUNT(starName) FROM StarsIn;

SELECT COUNT(DISTINCT starName) FROM StarsIn;
```

The execution results are shown in Figure 2.



```
SELECT COUNT(DISTINCT starName) FROM StarsIn;
      avg
-----
89272727.272727272727
(1 row)

      sum
-----
982000000
(1 row)

      min
-----
20000000
(1 row)

      max
-----
200000000
(1 row)

      count
-----
7
(1 row)

      count
-----
7
(1 row)

      count
-----
6
(1 row)

postgres=#
```

Figure 2: Aggregation Operators

1.3 Grouping

```
SELECT studioName, SUM(length) FROM Movies
GROUP BY studioName;

SELECT studioName, SUM(length), Count(length), AVG(length) FROM Movies
GROUP BY studioName;

SELECT studioName, Count(*) FROM Movies
GROUP BY studioName;

SELECT studioName FROM Movies
GROUP BY studioName;
```

The execution results are shown in Figure 3.

```

GROUP BY name;

SELECT name, Count(*)
FROM MovieExec,
      Movies
WHERE producerC = cert
GROUP BY name;
-----
name | sum
-----
Bryan Singer | 106
George Roy Hill | 187
J.J. Abrams | 348
AAA | 134
Garry Marshall | 119
Calvin Coolidge | 445
Dino De Laurentiis | 134
George Lucas | 235
(8 rows)

name | count
-----
Bryan Singer | 1
George Roy Hill | 1
J.J. Abrams | 3
AAA | 1
Garry Marshall | 1
Calvin Coolidge | 4
Dino De Laurentiis | 1
George Lucas | 2
(8 rows)

postgres=#
+ Shell Shell No. 2 Shell No. 3
~: psql

```

Figure 4: Grouping with More Than One Relation

```

GROUP BY studioName;
-----
studioName | sum
-----
MGM | 476
Universal | 287
Fox | 235
Disney | 119
Paramount | 477
(5 rows)

studioName | sum | count | avg
-----
MGM | 476 | 3 | 158.66666666666667
Universal | 287 | 2 | 143.50000000000000
Fox | 235 | 2 | 117.50000000000000
Disney | 119 | 1 | 119.00000000000000
Paramount | 477 | 4 | 119.25000000000000
(5 rows)

studioName | count
-----
MGM | 4
Universal | 2
Fox | 2
Disney | 1
Paramount | 4
(5 rows)

studioName
-----
MGM
Universal
Fox
Disney
Paramount
(5 rows)

postgres=#
+ Shell Shell No. 2 Shell No. 3
~: psql

```

Figure 3: Grouping

1.4 Grouping with More Than One Relation

```

SELECT name, SUM(length)
FROM MovieExec,
      Movies
WHERE producerC = cert
GROUP BY name;

SELECT name, Count(*)
FROM MovieExec,
      Movies
WHERE producerC = cert
GROUP BY name;

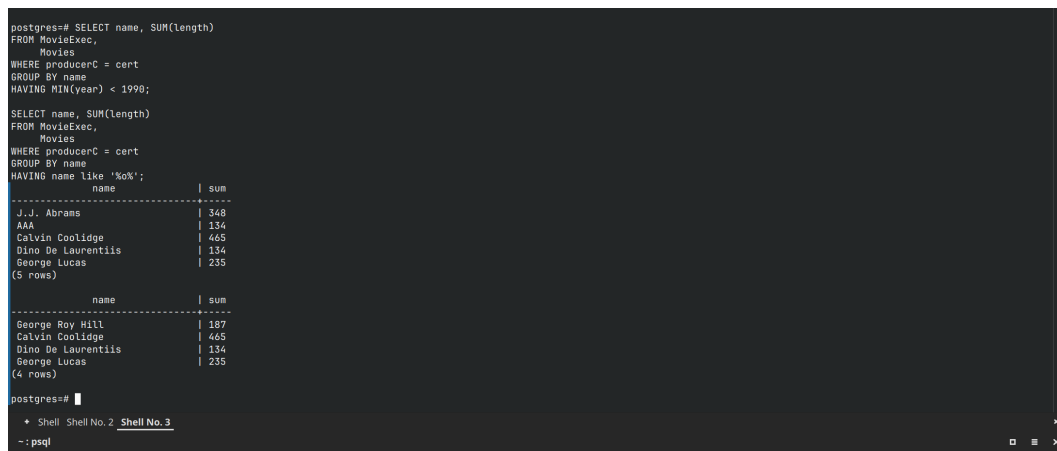
```

1.5 Having Clause

```
SELECT name, SUM(length)
FROM MovieExec,
     Movies
WHERE producerC = cert
GROUP BY name
HAVING MIN(year) < 1990;

SELECT name, SUM(length)
FROM MovieExec,
     Movies
WHERE producerC = cert
GROUP BY name
HAVING name like '%o%';
```

The execution results are shown in Figure 5.



```
postgres=# SELECT name, SUM(length)
FROM MovieExec,
     Movies
WHERE producerC = cert
GROUP BY name
HAVING MIN(year) < 1990;

SELECT name, SUM(length)
FROM MovieExec,
     Movies
WHERE producerC = cert
GROUP BY name
HAVING name like '%o%';
```

name	sum
J.J. Abrams	1348
AAA	134
Calvin Coolidge	465
Dino De Laurentiis	134
George Lucas	235

(5 rows)

name	sum
George Roy Hill	187
Calvin Coolidge	465
Dino De Laurentiis	134
George Lucas	235

(4 rows)

```
postgres=#
```

Figure 5: Having Clause

1.6 Insertion with Subqueries

```
INSERT INTO Studio(name)
SELECT DISTINCT studioName
FROM Movies
WHERE studioName NOT IN
     (SELECT name
      FROM Studio);
```

The execution results are shown in Figure 6.

```

      name      | sum
-----+-----
George Roy Hill | 187
Calvin Coolidge | 465
Dino De Laurentiis | 134
George Lucas    | 235
(4 rows)

postgres=# INSERT INTO Studio(name)
SELECT DISTINCT studioName
FROM Movies
WHERE studioName NOT IN
      (SELECT name
      FROM Studio);
INSERT 0 0
postgres=#

```

Figure 6: Insertion with Subqueries

1.7 Deletion

```

DELETE
FROM StarsIn
WHERE movieTitle = 'The Maltese Falcon'
      AND movieYear = 1942
      AND starName = 'Sydney Greenstreet';

```

the execution results are shown in Figure 7.

```

postgres=# INSERT INTO Studio(name)
SELECT DISTINCT studioName
FROM Movies
WHERE studioName NOT IN
      (SELECT name
      FROM Studio);
INSERT 0 0
postgres=# SET SQL_SAFE_UPDATES = 0;
postgres=# DELETE
FROM StarsIn
WHERE movieTitle = 'The Maltese Falcon'
      AND movieYear = 1942
      AND starName = 'Sydney Greenstreet';
ERROR: unrecognized configuration parameter "sql_safe_updates"
DELETE 0
postgres=#

```

Figure 7: Deletion

1.8 Update

```

UPDATE MovieExec
SET name = CONCAT('Pres. ', name)
WHERE cert IN (SELECT presC FROM Studio);

```

the execution results are shown in Figure 8.

```

      AND starName = 'Sydney Greenstreet';
ERROR: unrecognized configuration parameter "sql_safe_updates"
DELETE 0
postgres=# UPDATE MovieExec
SET name = CONCAT('Pres. ', name)
WHERE cert IN (SELECT presC FROM Studio);
UPDATE 5
postgres=#

```

Figure 8: Update

2 Exercise 6.4.1

Write each of the queries in Exercise 2.4.1 in SQL, making sure that duplicates are eliminated.

- `Product(maker, model, type)`
- `PC(model, speed, ram, hd, price)`
- `Laptop(model, speed, ram, hd, screen, price)`
- `Printer(model, color, type, price)`

- What PC models have a speed of at least 3.00?
- Which manufacturers make laptops with a hard disk of at least 100GB?
- Find the model number and price of all products (of any type) made by manufacturer B.
- Find the model numbers of all color laser printers.
- Find those manufacturers that sell Laptops, but not PC's.
- Find those hard-disk sizes that occur in two or more PC's.

2.1 Solutions

```
-- (a)
SELECT model
FROM PC
WHERE speed >= 3.00;

-- (b)
SELECT DISTINCT maker
FROM Product
JOIN Laptop ON Product.model = Laptop.model
WHERE hd >= 100;

-- (c)
SELECT model, price
FROM Product
WHERE maker = 'B';

-- (d)
SELECT model
FROM Printer
WHERE color = true AND type = 'laser';
```

```
-- (e)
SELECT DISTINCT maker
FROM Product
WHERE type = 'laptop'
      AND maker NOT IN (SELECT maker FROM Product WHERE type = 'pc');

-- (f)
SELECT hd
FROM PC
GROUP BY hd
HAVING COUNT(*) >= 2;
```

3 Exercise 6.4.6(a-i)

Write the following queries, based on the database schema

The database schema is:

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)

of Exercise 2.4.1, and evaluate your queries using the data of that exercise.

- Find the average speed of PC's.
- Find the average speed of laptops costing over \$1000.
- Find the average price of PC's made by manufacturer "A."
- Find the average price of PC's and laptops made by manufacturer "D."
- Find, for each different speed, the average price of a PC.
- Find for each manufacturer, the average screen size of its laptops.
- Find the manufacturers that make at least three different models of PC.
- Find for each manufacturer who sells PC's the maximum price of a PC.
- Find, for each speed of PC above 2.0, the average price.

3.1 Solutions

```
-- (a)
SELECT AVG(speed) FROM PC;

-- (b)
SELECT AVG(speed) FROM Laptop WHERE price > 1000;

-- (c)
SELECT AVG(price) FROM PC WHERE model IN (SELECT model FROM Product
    WHERE maker = 'A');

-- (d)
SELECT AVG(price) FROM (
    SELECT price FROM PC WHERE model IN (SELECT model FROM Product WHERE
        maker = 'D')
    UNION ALL
    SELECT price FROM Laptop WHERE model IN (SELECT model FROM Product
        WHERE maker = 'D')
) AS Prices;

-- (e)
SELECT speed, AVG(price) FROM PC GROUP BY speed;

-- (f)
SELECT maker, AVG(screen) FROM Product JOIN Laptop ON Product.model =
    Laptop.model GROUP BY maker;

-- (g)
SELECT maker FROM Product WHERE type = 'pc' GROUP BY maker HAVING COUNT(
    DISTINCT model) >= 3;

-- (h)
SELECT maker, MAX(price) FROM Product JOIN PC ON Product.model = PC.
    model GROUP BY maker;

-- (i)
SELECT speed, AVG(price) FROM PC WHERE speed > 2.0 GROUP BY speed;
```

4 Excerise 6.3.9

Using the two relations

- Classes(class, type, country, numGuns, bore, displacement)

- Ships(name, class, launched)

from our database schema of Exercise 2.4.3, write a SQL query that will produce all available information about ships, including that information available in the C:Classes relation. You need not produce information about classes if there are no ships of that class mentioned in Ships.

4.1 Solutions

```
SELECT
    Ships.name,
    Ships.class,
    Ships.launched,
    Classes.type,
    Classes.country,
    Classes.numGuns,
    Classes.bore,
    Classes.displacement
FROM Ships
LEFT JOIN Classes ON Ships.class = Classes.class;
```